

Lecture 13: SLAM I: Non-parametric Localization and EKF SLAM

13.1 Introduction

This lecture continues the discussion of localization with non-parametric localization. Afterwards, we transition into the topic of Simultaneous Localization and Mapping (SLAM).

Section 13.2 discusses non-parametric localization, which is more apt for problem setups when the Gaussian assumptions for the probability distributions no-longer hold. Overall, the non-parametric approach is quite similar to the parametric solution, but uses a histogram of particles to represent the belief instead. Specifically, we will be covering Monte Carlo Localization (MCL) as an example of non-parametric localization.

Section 13.3 begins our discussion on SLAM, an algorithm to both find the path for a robot and acquire a map of an unknown environment. As the name implies, SLAM algorithms include both the robot's pose AND a map of the environment in the state. In this lecture, we will cover EKF SLAM, a parametric SLAM approach that extends from EKF localization. EKF SLAM maintains a Gaussian assumption for the approximation of belief and uses a feature-based map to identify the environment.

13.2 Non-parametric Localization

So far, we have seen localization based on parametric filtering techniques. Now we will look at utilizing some of the non-parametric filtering techniques discussed earlier to formulate non-parametric localization. The following two subsections discuss Grid Localization and Monte Carlo Localization.

13.2.1 Grid Localization

Grid localization is a variant of the discrete Bayes filter and is mainly based on the principles of histogram filtering. The main idea is to approximate the posterior belief using a histogram filter over the grid decomposition of the state space:

$$\text{dom}(X_t) = x_{1,t} \cup x_{2,t} \cup \dots x_{K,t}$$

Hence, at time t , the state space is partitioned into K grids. Following this, the posterior is represented as a collection of discrete probability values:

$$\text{bel}(x_t) = \{p_{k,t}\}$$

where $p_{k,t}$ represents the probability of belonging to grid cell k at time t . Similar to the histogram filter, the motion and measurement models are also discretized with an additional conditioning on the map m . The

pseudocode in Figure 13.1 follows the same pattern as earlier localization methods. The input consists of the belief at time $t - 1$, the control action u_t , the measurement z_t , and the map m of the environment. The algorithm outputs the updated belief at time t .

The algorithm has a prediction step where the predicted belief $\{\bar{p}_{k,t}\}$ is calculated using a total probability theorem as shown in the algorithm. The total probability theorem uses the probability belief calculated for each cell i at time $t - 1$ defined as $p_{i,t-1}$ and the probability of reaching cell k and the map m . The last step normalizes the probability using the factor η , which is computed along with the probability belief $p_{k,t}$.

```

Data:  $\{p_{k,t-1}\}, u_t, z_t, m$ 
Result:  $\{p_{k,t}\}$ 
 $\eta \leftarrow 0$ 
foreach  $k$  do
   $\bar{p}_{k,t} = \sum_i \bar{\eta} |x_{k,t}| p(\hat{x}_k | u_t, \hat{x}_i, m) p_{i,t-1};$ 
end
foreach  $k$  do
   $p_{k,t} = p(z_t | \hat{x}_k, m) \bar{p}_{k,t};$ 
   $\eta \leftarrow \eta + p_{k,t};$ 
end
foreach  $k$  do
   $p_{k,t} = \eta^{-1} p_{k,t};$ 
end
Return  $\{p_{k,t}\}$ 

```

Figure 13.1: Pseudocode for Grid Localization

To decompose the state space into a discrete set of cells, there are multiple methods. One possible choice is topological, where decomposition is based on significant spaces. Another method is a uniform decomposition of state space into cells of equal size. Grid localization has a trade-off between retaining information from the continuous state space and complexity of computation. As the system uses more grids with higher resolution, the computational complexity increases. Some of the techniques used to reduce the computation complexity are pre-caching, sensor subsampling, delayed motion updates and selective updating.

13.2.2 Monte Carlo Localization

Monte Carlo Localization is a localization algorithm that follows from the particle filtering technique. For this algorithm, the belief at time t is defined as:

$$bel(x_t) = \chi = \{x_t^{[1]}, x_t^{[2]}, \dots, x_t^{[M]}\}$$

where $x_t^{[i]}$ is the i -th particle at time t . As in grid localization, the input consists of the control action u_t , the measurement z_t , and the map m of the environment along with the belief at time $t - 1$. The output is a new set of particles representing the belief at time t .

During the prediction step, M particles $\{x_t^{[m]}\}_{m=1}^M$ are sampled. Precisely, the particle $x_t^{[i]}$ is sampled from the motion model conditioned on the current particle $x_{t-1}^{[m]}$ and the control action u_t . A weight $w_t^{[m]}$ is assigned to each particle which indicates the confidence in that particle. This weight is calculated from the measurement model conditioned on the measurement z_t , the sampled particle $x_t^{[m]}$ and the map m .

The predicted belief $\bar{\chi}_t$ is a set containing tuples of particles and their associated weights. Then, in order to update χ_t , the algorithm randomly samples M particles with a probability distribution based on the confidence weights. So a particle with a higher weight could appear multiple times in the updated set χ_t , whereas a particle with a weight set to 0 would never appear. Thus, through this predicting and updating process, the belief narrows down to a set of particles with a high level of confidence. The pseudo-code for the algorithm is shown in Figure 13.2.

```

Data:  $\mathcal{X}_{t-1}, u_t, z_t, \mathbf{m}$ 
Result:  $\mathcal{X}_t$ 
 $\bar{\mathcal{X}}_t = \mathcal{X}_t = \emptyset;$ 
for  $i = 1$  to  $M$  do
    Sample  $x_t^{[m]} \sim p(x_t | u_t, x_{t-1}^{[m]}, \mathbf{m});$ 
     $w_t^{[m]} = p(z_t | x_t^{[m]}, \mathbf{m});$ 
     $\bar{\mathcal{X}}_t = \bar{\mathcal{X}}_t \cup (x_t^{[m]}, w_t^{[m]});$ 
end
for  $i = 1$  to  $M$  do
    Draw  $i$  with probability  $\propto w_t^{[i]};$ 
    Add  $x_t^{[i]}$  to  $\mathcal{X}_t;$ 
end
Return  $\mathcal{X}_t$ 

```

Figure 13.2: Pseudocode for Monte Carlo Localization [1]

Figure 13.3 provides an illustration of the Monte Carlo Localization algorithm. At the beginning, the robot is unaware of its location and the particles (black dots) are spread out widely in block (a). As the updating of particles progresses in block (b), the robot is able to narrow down its belief on its position to fewer regions shown by the clusters of particles. Eventually, the robot is very confident about its position as shown by the single cluster of particles in block (c). This algorithm can address some cases such as global localization, robot kidnapping, position tracking, and dynamic environments.




Figure 13.3: Illustration of Monte Carlo Localization [1]




Figure 13.4: Difference between online and full SLAM

13.3 Simultaneous Localization and Mapping (SLAM)

13.3.1 Motivations and general description

SLAM is the computation problem of constructing and updating a map of an unknown environment while simultaneously keeping track of an agent's location within it, using measurements $z_{1:t}$ and the controls $u_{1:t}$ up to time t . This data is gathered by the robot's proprioceptive and exteroceptive sensors. SLAM is difficult because both the estimated path and extracted features are corrupted by noise. We assume that the initial location of the robot is known and the others are not.

According to this terminology, the full SLAM problem consists of estimating the joint posterior probability of the robot path $x_{1:t}$ and the map m , from the odometry $u_{1:t}$ and observations $z_{1:t}$

$$p(x_{1:t}, m | z_{1:t}, u_{1:t})$$

The online SLAM problem aims to estimate the joint posterior over the robot's current pose x_t and the map m , from the data $z_{1:t}, u_{1:t}$

$$p(x_t, m | z_{1:t}, u_{1:t})$$

This is represented graphically in Figure 13.4.

13.3.2 Extended Kalman Filter localization with map uncertainty: EKF SLAM

EKF SLAM shares some similarities with EKF localization. First, we assume the existence of a feature-based map

$$m = \{m_1, \dots, m_N\}$$

where m_i is the i -th landmark. We denote the global frame coordinates of the i -th point landmark m_i with $(m_{i,x}, m_{i,y})$. Moreover, we assume the availability of a sensor that can measure the range and bearing pair $z = (r, \phi)$ of a landmark relative to the robot's local coordinate frame. Furthermore, in order to keep track of the uncertainty over the robot's pose x_t and the map m , we use Markov probabilistic models for motion, $p(x_t | x_{t-1}, u_{t-1})$, and perception, $p(z_t | x_t, m)$. These models, as well as our initial beliefs over the robot's pose and the map are assumed to be Gaussian.

Similarly to EKF localization, EKF SLAM aims to estimate the robot pose $x_t = (x, y, \theta)$ at time t with the following procedure:

- A set $z_t = \{z_t^1, z_t^2, \dots\}$ of range and bearing pairs $z_t^i = (r_t^i, \phi_t^i)$ is measured by the robot through its sensor.
- A landmark m_j is associated with each measurement z_t^i . We define the correspondence variable $c_t^i \in \{1, \dots, N+1\}$, such that $c_t^i = j$ if the i -th measurement z_t^i corresponds to the j -th landmark, and $c_t^i = N+1$ if no landmark corresponds to the i -th measurement.
- Given those new measurements and associations, the resulting posterior distribution on the robot's pose is computed.

We distinguish between two versions of EKF SLAM:

- The correspondence variables c_t^i are known, i.e.: given a measurement, we know the corresponding landmark (idealistic case).
- The correspondence variables c_t^i and the number of landmarks N are unknown. The algorithm must account for this uncertainty (usual case).

Unlike EKF localization, EKF SLAM does not assume knowledge of the landmarks' coordinates $\{(m_{i,x}, m_{i,y})\}_{i=1}^N$. Hence, EKF SLAM estimates the coordinates of all landmarks.

13.3.3 EKF SLAM with known correspondences

In the case of known correspondences, the approach to track the uncertainty over both the robot's pose x_t and the map m is to consider an augmented state $y_t \in \mathbb{R}^{3+2N}$ defined as:

$$y_t := \begin{pmatrix} x_t \\ m \end{pmatrix} = (x, y, \theta, m_{1,x}, m_{1,y}, \dots, m_{N,x}, m_{N,y})^T \quad (13.1)$$

Hence, the augmented state y_t comprises the robot's pose x_t and the representation of the map m through its landmarks' coordinates in the global coordinate frame. At time t , given all the measurements $z_{1:t}$ and all the inputs $u_{1:t}$ up to time t , the goal is to calculate the online posterior distribution over the augmented state y_t :

$$p(y_t | z_{1:t}, u_{1:t})$$

Note that here we do not compute a posterior over the variables c_t , since we assume known correspondences.

To apply the EKF machinery, we treat SLAM as an estimation problem and assume all landmarks are static. In order to cast everything into an EKF filtering problem, we need to specify motion and sensing models. Assume linear motion model with state $x_t = (x, y, \theta)$:

$$y_t = g(u_t, y_{t-1}) + \epsilon_t, \quad \epsilon_t \sim N(0, R_t), \quad G_t := J_g(u_t, y_{t-1}) \quad (13.2)$$

where $g(u_t, y_{t-1})$ is the dynamics of our entire state space, including the pose of the robot and the position of the landmarks. Hence, g is a $3 + 2N$ vector whose first three components describe the dynamics of the robot and the rest $2N$ are all zeros since the landmarks are static. The covariances between the $3 + 2N$ variables on system dynamics are stored in the covariance matrix, R_t . For the noise model of our

dynamics, R_t has dimension $3 + 2N$ by $3 + 2N$, and, for similar reason, only the upper left 3 by 3 sub-matrix representing the noise in robot pose is nonzero. The matrix G_t is the Jacobian of g at (u_t, y_{t-1}) .

We instantiate the measurement model in the same way as for EKF localization and assume a range and bearing measurement model:

$$\hat{z}_t^i = \begin{pmatrix} \sqrt{(m_{j,x} - x)^2 + (m_{j,y} - y)^2} \\ \text{atan2}(m_{j,y} - y, m_{j,x} - x) - \theta \end{pmatrix} + \delta_t := h(y_t, j) + \delta_t \quad (13.3)$$

where

$$\delta_t \sim N(0, Q_t), \quad Q_t = \begin{pmatrix} \sigma^2 & 0 \\ 0 & \sigma_\phi^2 \end{pmatrix} \text{ and } c_t^i = j$$

The function that maps the states into the measurement z_t^i is given by range and bearing, so this is identical to the case of EKF localization. The only difference is that we are mapping a much larger state vector with some noise into a 2D vector. As it was the case for the motion model, we approximate the measurement model according to a linear function by taking a Taylor expansion around the mean of the predicted belief over y_t .

$$h(y_t, j) \approx h(\bar{\mu}_t, y) + H_t^i(y_t - \bar{\mu}_t) \quad (13.4)$$

where

$$H_t^i := \frac{\partial h(\bar{\mu}_t, j)}{\partial y_t}$$

When taking the derivative of the function h to obtain the Jacobian, we will have zero terms because we are taking derivatives with respect to state variables that do not impact the measurement model. Then, the derivative of h with respect to x_t and m_j can be expressed as

$$h_t^i = \frac{\partial h(\bar{\mu}_t, j)}{\partial(x_t, m_j)} = \begin{pmatrix} \frac{\bar{\mu}_{t,x} - \bar{\mu}_{j,x}}{\sqrt{q_{t,j}}}, \frac{\bar{\mu}_{t,y} - \bar{\mu}_{j,y}}{\sqrt{q_{t,j}}}, 0, \frac{\bar{\mu}_{j,x} - \bar{\mu}_{t,x}}{\sqrt{q_{t,j}}}, \frac{\bar{\mu}_{j,y} - \bar{\mu}_{t,y}}{\sqrt{q_{t,j}}} \\ \frac{\bar{\mu}_{j,y} - \bar{\mu}_{t,y}}{q_{t,j}}, \frac{\bar{\mu}_{j,x} - \bar{\mu}_{t,x}}{q_{t,j}}, -1, \frac{\bar{\mu}_{t,y} - \bar{\mu}_{j,y}}{q_{t,j}}, \frac{\bar{\mu}_{t,x} - \bar{\mu}_{j,x}}{q_{t,j}} \end{pmatrix} \quad (13.5)$$

where $\bar{\mu}_{t,x}$ and $\bar{\mu}_{t,y}$ are the current estimations of the x and y coordinates of the robots, and $\bar{\mu}_{j,x}$ and $\bar{\mu}_{j,y}$ are the expected values of x and y coordinates with respect to the j^{th} landmark. And

$$q_j^t := (\bar{\mu}_{j,x} - \bar{\mu}_{t,x})^2 + (\bar{\mu}_{j,y} - \bar{\mu}_{t,y})^2.$$

Define $F \in \mathbb{R}^{5 \times (3+2N)}$ as

$$F_j^i = \begin{pmatrix} 1 & 0 & 0 & 0 \cdots 0 & 0 & 0 & 0 \cdots 0 \\ 0 & 1 & 0 & 0 \cdots 0 & 0 & 0 & 0 \cdots 0 \\ 0 & 0 & 1 & 0 \cdots 0 & 0 & 0 & 0 \cdots 0 \\ 0 & 0 & 0 & 0 \cdots 0 & 1 & 0 & 0 \cdots 0 \\ 0 & 0 & 0 & \underbrace{0 \cdots 0}_{2j-2} & 0 & 1 & \underbrace{0 \cdots 0}_{2N-2j} \end{pmatrix} \quad (13.6)$$

Then, we have the following factorization:

$$H_t^i = h_t^i F_{x,j} \quad (13.7)$$

The following step is initializing the algorithm. Usually the initial pose is taken to be the origin of the coordinate system. In this case, the map is unknown. The initial belief is then expressed as

$$\mu_0 = (0, 0, \dots, 0)^T \quad (13.8)$$

Although an arbitrary choice, taking the initial belief as a zero vector can greatly reduce computation in later operations. The next step is initializing the covariance matrix for augmented state vectors. Assuming that the robot has known pose and is at the origin of an unknown map, we have zero covariances for the top left 3*3 block of the matrix as initialization for pose variables. Since the locations of landmarks are unknown, the covariances of landmarks with respect to themselves are taken to be infinite, which is represented by large numbers in practice. In addition, there is no information on the cross correlations between different landmarks, and hence the off-diagonal elements in the covariance matrix are initialized as zeros.

$$\Sigma_0 = \begin{pmatrix} 0 & 0 & 0 & 0 & \cdots & 0 \\ 0 & 0 & 0 & 0 & \cdots & 0 \\ 0 & 0 & 0 & 0 & \cdots & 0 \\ 0 & 0 & 0 & \infty & \cdots & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \infty \end{pmatrix} \quad (13.9)$$

When the robot observes a landmark for the first time, it initializes the position for that landmark based units estimated current position obtained using range and bearing measurement model. Hence the landmark estimates $(\bar{\mu}_{j,x}, \bar{\mu}_{j,y})^T$ can be determined as:

$$\begin{pmatrix} \bar{\mu}_{j,x} \\ \bar{\mu}_{j,y} \end{pmatrix} = \begin{pmatrix} \bar{\mu}_{t,x} \\ \bar{\mu}_{t,y} \end{pmatrix} + \begin{pmatrix} r_t^i \cos(\phi_t^i + \bar{\mu}_{t,\theta}) \\ r_t^i \sin(\phi_t^i + \bar{\mu}_{t,\theta}) \end{pmatrix} \quad (13.10)$$

The first term on the right-hand side is the estimated position of the robot. In the second term, r_t^i , represents the distance between robot and landmark, and the trigonometry terms on the relative angles stores the directional information. EKF SLAM algorithm, shown in Figure (13.5), is very similar to the EKF localization algorithm. The biggest difference is that new landmarks may be discovered and needs to be accounted for. Other main differences are augmented state vector, augmented dynamics (with trivial dynamics for the landmarks), and augmented measurement Jacobian. Also similar to EKF localization, this algorithm contains a Bayesian filter and the prediction step propagates the expected value of our state factor, which is an augmented state vector, and the covariance of our predicted belief using the motion model. Since landmarks are assumed to be static, only the first three elements of the state vector and the top left 3 by 3 block of the covariance matrix are updated. Therefore, only the mean of the covariance belief is manipulated in the prediction step in accordance to the motion model. And the step only affects elements with beliefs distribution related to a robot pose. Afterwards, information about the measurements will be incorporated. Each measurement is processed sequentially as in EKF localization. With the assumption of known correspondence variables, once measurements z_t^i are obtained, we have correspondence variable that shows c_t^i measurement corresponding to the landmark j . If the landmark j has not been seen before, its position will be initialized. The rest of steps are the same as in EKF localization. Time and memory complexity for EKF SLAM algorithm are $O(N^2)$.

13.3.4 EKF SLAM with unknown correspondences

The general idea of EKF SLAM with unknown correspondences is to use maximum likelihood estimation to decide the measurement correspondences during the measurement step. The rest of the algorithm follows similar to the EKF SLAM algorithm assuming known correspondences discussed previously.

```

Data:  $(\mu_{t-1}, \Sigma_{t-1}), u_t, z_t, N_{t-1}$ 
Result:  $(\mu_t, \Sigma_t)$ 
 $N_t = N_{t-1};$ 
 $\bar{\mu}_t = g(u_t, \mu_{t-1});$ 
 $\bar{\Sigma}_t = G_t \Sigma_{t-1} G_t^T + R_t;$ 
foreach  $z_t^i = (r_t^i, \phi_t^i)^T$  do
     $\begin{pmatrix} \bar{\mu}_{N_t+1,x} \\ \bar{\mu}_{N_t+1,y} \end{pmatrix} = \begin{pmatrix} \bar{\mu}_{t,x} \\ \bar{\mu}_{t,y} \end{pmatrix} + \begin{pmatrix} r_t^i \cos(\phi_t^i + \bar{\mu}_{t,\theta}) \\ r_t^i \sin(\phi_t^i + \bar{\mu}_{t,\theta}) \end{pmatrix};$ 
    for  $k = 1$  to  $N_t + 1$  do
         $\hat{z}_t^k = \begin{pmatrix} \sqrt{(\bar{\mu}_{j,x} - \bar{\mu}_{t,x})^2 + (\bar{\mu}_{j,y} - \bar{\mu}_{t,y})^2} \\ \text{atan2}(\bar{\mu}_{j,y} - \bar{\mu}_{t,y}, \bar{\mu}_{j,x} - \bar{\mu}_{t,x}) - \bar{\mu}_{t,\theta} \end{pmatrix};$ 
         $H_t^k = h_t^k F_{x,k};$ 
         $S_t^k = H_t^k \bar{\Sigma}_t [H_t^k]^T + Q_t;$ 
         $\pi_k = (z_t^i - \hat{z}_t^k)^T [S_t^k]^{-1} (z_t^i - \hat{z}_t^k);$ 
    end
     $\pi_{N_t+1} = \alpha;$ 
     $j(i) = \text{argmin}_k \pi_k;$ 
     $N_t = \max\{N_t, j(i)\};$ 
     $K_t^i = \bar{\Sigma}_t [H_t^{j(i)}]^T [S_t^{j(i)}]^{-1};$ 
     $\bar{\mu}_t = \bar{\mu}_t + K_t^i (z_t^i - \hat{z}_t^{j(i)});$ 
     $\bar{\Sigma}_t = (I - K_t^i H_t^{j(i)}) \bar{\Sigma}_t;$ 
end
 $\mu_t = \bar{\mu}_t$  and  $\Sigma_t = \bar{\Sigma}_t;$ 
Return  $(\mu_t, \Sigma_t)$ 

```

Hypothesis
for new
landmark

Mahalanobis
distance

Hypothesis test

Figure 13.5: Pseudocode for EKF SLAM with unknown correspondences.

This algorithm's more general applicability comes at the cost of being less robust. The key difference in the interface of the algorithm is that, in addition to estimating the positions of each of the landmarks, it also may estimate the number of landmarks. Therefore we consider N_t , the estimated number of landmarks at time t to be an input. At the measurement step, we loop over all measurements z_t^i . Then, for each of the N_t landmarks, we calculate the likelihood of each z_t^i , corresponding to each landmark.

The first step is to form \hat{z}_t^k , the expected measurement at time t given our belief of landmark k , $\bar{\mu}_k$. In the case where we receive range and bearing measurements, this is:

$$\hat{z}_t^k = \begin{pmatrix} \sqrt{(\bar{\mu}_{k,x} - \bar{\mu}_{t,x})^2 + (\bar{\mu}_{k,y} - \bar{\mu}_{t,y})^2} \\ \text{atan2}(\bar{\mu}_{k,y} - \bar{\mu}_{t,y}, \bar{\mu}_{k,x} - \bar{\mu}_{t,x}) - \bar{\mu}_{t,\theta} \end{pmatrix} \quad (13.11)$$

We can then calculate the likelihood z_t^i corresponding to our landmark by taking the Mahalanobis distance between z_t^i and \hat{z}_t^k . The Mahalanobis distance corresponds to the z-score of the measurement z_t^i under our measurement model centered at \hat{z}_t^k . Using the notation similar to before, we form

$$H_t^k = h_t^k F_{x,k} \quad (13.12)$$

which is the Jacobian of the measurement model with respect to the k th landmark at time t . We then use this Jacobian to project the error covariance of the state at time t to the error covariance of our measure-

ment:

$$S_t^k = H_t^k \bar{\Sigma}_t [H_t^k]^T + Q_t \quad (13.13)$$

Finally, we calculate the Mahalanobis distance as:

$$\pi_k = (z_t^i - \hat{z}_t^k) [S_t^k]^{-1} (z_t^i - \hat{z}_t^k) \quad (13.14)$$

We can then estimate $j(i)$, the correspondence of measurement Z_t^i as:

$$j(i) = \underset{k \in 1, \dots, N_t+1}{\operatorname{argmin}} \pi_k \quad (13.15)$$

We also set a threshold distance $\pi_{N_t+1} = \alpha$. If it happens that α is the minimum distance between measurement z_t^i and each landmark, then we treat z_t^i as having come from a new landmark, $N_t + 1$, and we initialize its belief as the expected position given the measurement z_t^i :

$$\begin{pmatrix} \bar{\mu}_{N_t+1,x} \\ \bar{\mu}_{N_t+1,y} \end{pmatrix} = \begin{pmatrix} \bar{\mu}_{t,x} \\ \bar{\mu}_{t,y} \end{pmatrix} + \begin{pmatrix} r_t^i \cos(\phi_t^i + \bar{\mu}_{t,\theta}) \\ r_t^i \sin(\phi_t^i + \bar{\mu}_{t,\theta}) \end{pmatrix} \quad (13.16)$$

We would also increment N_t by 1. After we have formed the correspondence function $j(i)$, we can simply proceed with the measurement step as before, to update our beliefs. The full pseudo-code for the algorithm is shown in Figure 13.5.

As mentioned before, this algorithm's broader applicability comes at the cost of being less robust. In particular, extraneous measurements can result in the creation of fake landmarks, which will then propagate forward to future steps uncorrected. Furthermore, the EKF can diverge if non-linearities are large.

There are several techniques to mitigate these issues, such as:

- Outlier rejection schemes [3] employing provisional landmark lists, etc.
- Strategies to enhance the distinctiveness of landmarks which may involve prior knowledge or assumptions, such as spatial arrangements, signatures, geometric constraints, etc.

However, the central dilemma of EKF SLAM is that accurate localization typically requires dense maps, while EKF requires sparse maps due to the quadratic complexity of updates.

13.4 Further References

1. "A review of recent developments in Simultaneous Localization and Mapping" offers a good overview of SLAM and its future directions of research:
<https://ieeexplore.ieee.org/abstract/document/6038117>
2. "Simultaneous Localisation and Mapping (SLAM): Part I The Essential Algorithms" clearly explains the probabilistic approach of SLAM, an overview of how EKF SLAM works, and future research opportunities.
<https://ieeexplore.ieee.org/document/1638022>

Contributors

Winter 2019: Tariq Zahroof, Ashar Alam, Andre Cornman, Ryan Cohen, Rohun Kulkarni, Jerin Joseph, Thomas Trzpit

Winter 2018: Malavika Bindhi, Charles Preston Hale, Arthur Ji, Jonathan Lacotte, Hongyi Zhao, Jerry Zheng